

Digital Twins for Predictive Maintenance in Automotive Aftersales

Maksim Bogachenkov
Syntegon
Stuttgart

wi23024@lehre.dhbw-stuttgart.de

Oliver Feng
Trumpf
Ditzingen

wi23030@lehre.dhbw-stuttgart.de

Joel Yeraí Martínez Campo
BITBW
Stuttgart

wi23194@lehre.dhbw-stuttgart.de

Abstract—This paper examines the effect of digital-twin-based predictive maintenance on the automotive aftersales domain, focusing on the German passenger-car market. Combining a structured literature review with a model-based gap analysis, we compare an interval-based current state against a digital-twin-enabled target state along the full customer journey, then quantify the business potential using an explicit, transparent assumption framework anchored to published German market data. The results show that proactive, data-driven service recommendations raise the repair capture rate from 40 to 70 percent, increasing wear-part gross profit by approximately 55 EUR per vehicle per year and 276,000 EUR annually across a 5,000-vehicle fleet. Customer touchpoints grow from 1.5 to roughly 2.5 per year, each representing an additional advisory and upselling opportunity. The findings establish a transparent, contestable model of the potential that dealerships and OEMs can use to evaluate and plan the integration of digital-twin-based predictive maintenance into their aftersales operations.

Index Terms—Digital Twin, Predictive Maintenance, Industry 4.0, Machine Learning, IoT, Condition Monitoring

I. INTRODUCTION

Today’s vehicles are connected by default and continuously transmit operating data [1], [2], so the information needed to judge their technical condition largely exists. Yet aftersales remains reactive: even well-connected vehicles are serviced on fixed time or mileage intervals rather than on the actual condition of their components [3]. Wear parts—brake pads and discs, tyres, ignition parts, or wipers—are replaced only once a fault becomes visible, so customers run them to failure and the unplanned repair often goes to a cheaper third-party garage [4]. For the dealership, which owns the customer relationship, this means losing both the repair and one of only a few yearly customer contacts—each also a chance to advise, retain, and sell [5], [6].

Digital-twin-based predictive maintenance (PdM) can turn this data into foresight. A digital twin (DT) maintains a synchronised virtual representation of the vehicle [7], against which machine-learning models forecast the remaining useful life of components and recommend replacement before failure—keeping the customer at the dealership and enabling transparent, targeted upselling. While the building blocks of DT-based PdM are well studied in industrial and aerospace settings [8], [9], their integration into automotive aftersales remains underexplored.

This paper therefore asks: *what effect does the proactive use of digital twins and predictive maintenance have on the automotive aftersales domain?* Focusing on the German market and working with explicitly stated assumptions, we (i) clarify the interplay of DT and PdM in aftersales, (ii) compare the predictive target state with the interval-based current state to reveal its potential, and (iii) derive recommendations for the dealership. The paper follows the IEEE conference structure: theoretical background, methodology, gap analysis, discussion, and conclusion.

II. RELATED WORK AND THEORETICAL BACKGROUND

A. Digital Twins

The concept of the digital twin was first introduced by Grieves [7] in a white paper on manufacturing excellence, where it described a virtual replica of a physical product maintained in synchrony with its real-world counterpart throughout the entire asset lifecycle. Originally conceived as a tool for improving production quality in industrial manufacturing, the idea drew on advances in simulation, sensor technology, and computational modelling to propose that a persistent virtual representation of a physical object could serve as the basis for monitoring, analysis, and decision-making. Over the subsequent decade, the concept migrated from its manufacturing origins into aerospace, energy, and infrastructure domains, before emerging as a foundational concept in Industry 4.0 discourse more broadly [8], [10].

In the automotive domain, this general architecture maps directly onto the connected vehicle. The physical entity is the vehicle itself, equipped with on-board sensors and an On-Board Diagnostic (OBD) interface. The virtual representation is a cloud-hosted model mirroring component states, degradation histories, and operating conditions. The synchronising link is the OTA telematics connection, which transmits telemetry to the cloud and returns configuration updates to the vehicle [1], [2]. Figure ?? illustrates this three-layer architecture. Lopes et al. [1] position this configuration within a broader software ecosystem in which the digital twin acts as the authoritative source of vehicle state for all downstream services. Architecturally, automotive digital twins are typically realised as modular, microservice-based platforms. A data ingestion layer normalises heterogeneous sensor streams—including vibration, temperature, voltage, and Controller Area

Network (CAN)-bus signals — before forwarding them to a model layer that maintains the current state of each monitored component [11], [12]. The outputs are then exposed to external consumers such as Original Equipment Manufacturer (OEM) backends or customer-facing notification services [1], [9], forming the basis for predictive maintenance strategies.

B. Vehicle Maintenance Strategies

Vehicle maintenance strategies form a maturity hierarchy: reactive maintenance repairs only after failure [9]; preventive maintenance schedules service at fixed mileage or time intervals, leading to either premature replacement or undetected wear [9], [11]; condition-based maintenance (CBM) couples the service decision to sensor thresholds [11]; and PdM extends CBM by forecasting the future evolution of degradation, estimating the remaining time or distance until end-of-life is reached [9], [13].

PdM relies on three data layers: on-board sensors delivering high-frequency signals such as vibration, current, voltage, and cell temperature [11]–[13]; the OBD interface providing standardised diagnostic codes [2], [14]; and OTA telematics enabling continuous low-latency exchange with the backend [1], [2], [12]. Together, these layers transform the vehicle from a periodically inspected asset into a continuously observed system, a shift first documented in 2009 by Zhang et al. [14].

PdM at this scale is not a stand-alone technique but a service layer built on a DT of the vehicle. As Lopes et al. [1] and Hadraoui et al. [11] describe, the DT supplies the synchronised virtual representation against which sensor streams are interpreted, anomalies detected, and remaining-useful-life estimates computed. Without a DT, telemetry must be processed in isolation; without PdM, the DT remains a passive monitor. In current automotive DT architectures, predictive analytics is therefore typically positioned as a dedicated microservice within the DT's analytics core [1], [9], [12].

C. Machine Learning in PdM

Three problem classes structure machine-learning (ML)-based PdM. Classification assigns sensor patterns to discrete fault categories using Support Vector Machines, Random Forests, or Convolutional Neural Networks on time-frequency representations of vibration signals [9], [13]. Regression targets Remaining Useful Life (RUL) estimation, dominated by Long Short-Term Memory (LSTM) networks for their ability to model long-range temporal dependencies in degradation signals [9], [12]. Anomaly detection operates without explicit labels, learning nominal behaviour and flagging deviations as failure precursors [9], [12]. AbouGrad et al. report classification accuracies of approximately 89% for vibration-based fault detection and 92% for LSTM-based RUL estimation in automotive contexts, though they note that these figures depend strongly on dataset quality and operating conditions.

Data availability is a persistent constraint: AbouGrad et al. and Hadraoui et al. note that preprocessing consumes 70–90% of project time, genuine failure data is structurally scarce, and public benchmarks representative of in-service behaviour

are largely absent. Model-based and hybrid approaches that embed physical degradation laws into the learning pipeline are therefore preferred over purely data-driven ones, as they extrapolate more reliably under data scarcity [8], [11], [13]. Beyond accuracy, deployment in safety-relevant contexts raises the question of explainability: deep architectures achieve high accuracy but operate as opaque models whose decisions cannot readily be justified to engineers, customers, or regulators [9]. Explainable AI (XAI) techniques such as feature attribution and attention-based saliency are increasingly integrated into PdM pipelines to expose the reasoning behind a prediction, marking the conceptual bridge from smart to wise information systems [9].

D. Automotive Aftersales

Aftersales encompasses post-purchase activities that help customers access supplier services and resolve product issues through continuous maintenance and upkeep [15]. As a quality-focused product support activity, it creates a sustainable competitive advantage across industries, with automotive being a prime example [16], [17]. In the automotive context, it functions as a dedicated business module bundling maintenance, repair, and warranty offerings to maximise product use throughout the asset lifecycle while meeting financial and customer-performance targets [6], [18]. Building post-sales relationships strengthens customer loyalty, yields actionable insights, and creates upselling and cross-selling opportunities [19], making the automotive aftersales market increasingly attractive as demand for higher-quality service expands [6].

Historically, once the warranty expires, customer engagement with the dealership tends to decline. Within the automotive industry, aftersales services represent a particularly significant source of revenue, with profit margins considerably exceeding those generated through vehicle sales [20]. From an economic standpoint, the aftersales service market has been identified as substantially larger than the primary product sales market across various industries [10], [21], [22]. Retaining these customers requires better information systems and service processes [15]. There is growing recognition that proactively anticipating issues before they arise improves performance, efficiency, and operational excellence [23].

Proactive aftersales improves customer satisfaction and retention by resolving issues before customers encounter them, enabling more productive service interactions [5]. This requires tailoring services to customer needs through tight integration of operational and informational data [15]. ML and DT-based approaches enable predictive interventions and better asset lifecycle optimisation, delivering cost reduction, improved customer satisfaction, and extended asset lifetimes [24].

E. Synthesis and Research Gap

The preceding subsections establish three building blocks: the DT as a synchronised vehicle model, PdM as a prognostic analytics layer built on top of it, and aftersales as the customer-facing business context in which service decisions

are communicated and acted upon. Together, they form a potential end-to-end system in which the DT continuously feeds component-state data to ML-based PdM models, whose forecasts are surfaced through aftersales processes to enable proactive service appointments and retain customers at the dealership.

While the referenced contributions establish the technical foundations of DT-based PdM, existing work concentrates either on algorithmic aspects in isolation [9], [13] or on industrial applications [8]. The integration of DT-based PdM into automotive aftersales business models, particularly regarding its implications for customer interaction, remains underexplored. The present work addresses this gap.

III. METHODOLOGY

This paper combines a structured literature review with a model-based gap analysis to answer the research question. The two steps are designed to complement each other: the review establishes what is technically possible and identifies where existing work stops short; the gap analysis then quantifies the business consequences of that shortfall for a concrete reference case.

a) *Literature review.*: The theoretical grounding rests on a targeted review of recent academic literature on DT, PdM, and automotive aftersales. Sources were identified through database searches on IEEE Xplore, Scopus, and Google Scholar using term clusters centred on *digital twin*, *predictive maintenance*, *automotive*, and *aftersales*, supplemented by reference tracing from key contributions. Selection prioritised peer-reviewed work with quantitative or architectural results published from 2015 onward, reflecting the period in which automotive DT architectures reached practical maturity [1]; foundational conceptual works predating this window were included where they introduced concepts still in active use [7]. The resulting synthesis is presented in Section II.

b) *Gap analysis.*: The gap analysis proceeds in two stages. First, a lifecycle comparison (Section IV-A) contrasts the interval-based current state of aftersales against a DT/PdM-enabled target state along the customer journey, identifying the structural points at which reactive servicing causes value leakage. Second, a quantitative potential assessment translates the identified gap into economic KPIs using an explicit assumption framework. Because no public, in-service failure dataset exists for the German passenger-car market at the resolution required [9], the assessment is model-based: assumptions are anchored to published statistics where available and stated as set values otherwise, making the model transparent and contestable. A single economic lever—the repair capture rate c , the share of addressable wear-part work retained at the franchised dealership—is used to quantify the gap. The target capture rate is derived analytically from the mechanism assumptions rather than assumed directly, ensuring internal consistency. A one-way sensitivity analysis over c confirms that the economic case holds across a plausible parameter band [25].

IV. GAP ANALYSIS

A. Future and Present Lifecycle Comparison




Fig. 1. Aftersales lifecycle: current interval-based state (left) versus DT/PdM-enabled target state (right).

The target of this comparison is to expose the structural differences between two states of the aftersales lifecycle—one without and one with integration of digital-twin-based predictive maintenance—along the full customer journey from vehicle operation through service execution to post-repair follow-up. The comparison reveals at which points in the lifecycle reactive servicing causes value leakage and where a predictive architecture would redirect that value back to the franchised dealership.

a) *Current state.*: In the interval-based current state, the customer journey is driven by fixed manufacturer schedules and visible failure signals. The vehicle operates without continuous component-state monitoring; no wear forecast is computed during normal driving. Service appointments are initiated either by a time- or mileage-triggered reminder displayed on the instrument cluster, or by an acute fault manifesting as a warning lamp or perceptible degradation in braking or ride performance. Both triggers are late-stage events: the mileage interval is set conservatively to cover the worst-case fleet population, so most vehicles reach the workshop with substantial remaining part life; the fault-driven trigger often arrives after the wear limit has already been exceeded [9]. Because no proactive outreach precedes the fault, the customer exercises free choice of workshop at the moment of highest urgency, and the majority of these unplanned repairs are captured by independent third-party garages [26], [27]. The dealership's contact with the customer is therefore reduced to the statutory inspection cycle, limiting both retention and upselling opportunity.

b) *Target state.*: In the DT/PdM-enabled target state, the vehicle's digital twin continuously aggregates sensor streams, OBD signals, and OTA telemetry to maintain an up-to-date virtual representation of each monitored component [1]. Machine-learning models operating on this virtual representation compute remaining-useful-life estimates for wear parts such as brake pads and discs, tyres, and ignition components [12], [13]. When the forecast horizon falls below a configurable threshold, the system generates a proactive service recommendation communicated to the customer through

the OEM’s connected-car interface or the dealership’s CRM before any visible fault has occurred [5], [24]. The appointment is scheduled at the customer’s convenience, the required parts are pre-ordered, and the workshop interaction is transformed from an unplanned, fault-driven event into a planned, evidence-based service visit. Post-service, the digital twin resets the component baseline and resumes monitoring, closing the lifecycle loop. The dealership thereby intercepts wear-part replacements that would otherwise be lost to the independent aftermarket and gains structured contact points along the customer journey that the current state does not provide.

B. Assumption Framework

Because public, in-service failure datasets representative of this setting are largely absent [9], the comparison is built on an explicit assumption framework rather than on a proprietary dataset. Assumptions are divided into *context* assumptions, which fix the reference case (Table I), *wear-part* assumptions, which size the addressable work (Table II), and *mechanism* assumptions, which govern the transition from the current to the target state (Table III). Each is anchored to a published figure where one exists and flagged as a set value otherwise.

TABLE I
CONTEXT ASSUMPTIONS.

ID	Assumption	Value
A1	Reference vehicle	Mid-segment ICE passenger car, German market
A2	Annual mileage	12,300 km/yr [28]
A3	First-owner horizon	8 years [27]
A4	Workshop visits	1.5/yr [26]
A5	Considered wear parts	Brake pads/discs, tyres, wipers, 12 V battery, spark plugs
A6	Service gross margin	50 % (range 45–55 % [27])
A7	Franchised share	43 % [26]

The wear parts in A5 are deliberately limited to components whose wear is gradual, usage-driven and not readily perceived by the driver — the class for which a forecast adds the most decision value and which dominates both between-interval failures and the most common defect categories in the mandatory German vehicle safety inspection (Hauptuntersuchung, HU: brakes, lighting, tyres) [29]. Prescriptive maintenance and powertrain overhauls are out of scope. Typical replacement intervals and representative German market prices (parts plus labour, incl. VAT) are given in Table II; prices are published market ranges [30] and major automotive service providers used at their midpoints. Where a published range exists, a representative value within it is used; individual prices and intervals may fall toward either end—premium tyres raise the tyre figure, a conventional (non start–stop) battery lowers the battery figure—but because every KPI scales linearly with V_{wear} , the sensitivity analysis below brackets this residual uncertainty and no single assumption drives the result.

For the mechanism, we model a single lever: the *repair capture rate* c , the share of addressable wear-part work performed

TABLE II
WEAR-PART ASSUMPTIONS (MIDPOINTS). PRICES COVER PARTS PLUS LABOUR, INCL. VAT; MIDPOINTS OF PUBLISHED GERMAN MARKET RANGES [30].

Part (A5)	Interval	Events/yr (at A2)	Price
Brake pads (per axle)	40,000 km	0.31	€ 250
Brake discs (+pads, front)	120,000 km	0.10	€ 450
Tyres (set of 4)	45,000 km	0.27	€ 600
Wiper blades (pair)	1 year	1.00	€ 25
12 V battery (fitted)	6 years	0.17	€ 150
Spark plugs (set)	60,000 km	0.21	€ 150

at the franchised dealership rather than lost to an independent garage. Its complement $(1 - c)$ is the leakage. The current capture rate, the forecast accuracy, and the proactive-offer conversion are stated in Table III; the target capture rate is not assumed independently but derived from them.

The target capture rate is derived, not assumed: of the leaked share $(1 - c_{\text{cur}})$, the fraction that PdM forecasts (a) and that converts on a proactive offer (k) is recaptured. We set $a=1$ because customers cannot distinguish a true-positive forecast from a false-positive recommendation—every offer is evaluated and accepted or declined at rate $a \cdot k$ regardless of whether the underlying wear has fully materialised:

$$\begin{aligned}
 c_{\text{tgt}} &= c_{\text{cur}} + (1 - c_{\text{cur}}) \cdot a \cdot k \\
 &= 0.40 + 0.60 \cdot 1.0 \cdot 0.50 \\
 &= 0.40 + 0.30 = 0.70.
 \end{aligned} \tag{1}$$

TABLE III
MECHANISM ASSUMPTIONS (CURRENT STATE → TARGET STATE).

ID	Assumption (Symbol)	Value	Basis
A8	Current capture rate (c_{cur})	0.40	<43 % share [26]
A9	PdM forecast accuracy (a)	0.90	89–92 % [9]; $a=1$ in model
A10	Proactive-offer conversion (k)	0.50	Conserv. est. [5], [17]
—	Target capture rate (c_{tgt} , deriv.)	0.70	$c_{\text{cur}} + (1 - c_{\text{cur}}) \cdot a \cdot k$, $a=1$

This construction keeps the estimate conservative and never approaches full capture. The conversion rate ($k = 0.50$, A10) is a conservative estimate: proactive, personalised maintenance contact from the customer’s own dealership is consistent with findings that proactive contact aids uptake and loyalty [5], [17], yet price competition and the availability of independent garages prevent full conversion.

C. Potential Assessment

The calculation proceeds in three steps: the addressable wear-part value per vehicle, the captured value under each state, and the resulting KPIs at vehicle, fleet and lifecycle level.

Step 1 — Addressable wear-part value: For each part i , the expected annual revenue is its replacement rate λ_i [yr^{-1}] multiplied by its unit price p_i ; both are from Table II, with λ_i

computed on unrounded mileage (A2). Summing across A5 yields the annual addressable wear-part revenue per vehicle:

$$V_{\text{wear}} = \sum_i \lambda_i \cdot p_i$$

$$= 76.9 + 46.1 + 164.0 + 25.0 + 25.0 + 30.7$$

$$\approx \text{€ } 368 \text{ per vehicle per year.}$$

Step 2 — Captured value: Applying the capture rates from Table III:

$$\text{Captured}_{\text{cur}} = c_{\text{cur}} \cdot V_{\text{wear}} = 0.40 \cdot 368 \approx \text{€ } 147 \text{ /veh/yr,} \quad (2)$$

$$\text{Captured}_{\text{tgt}} = c_{\text{tgt}} \cdot V_{\text{wear}} = 0.70 \cdot 368 \approx \text{€ } 258 \text{ /veh/yr.} \quad (3)$$

Applying the service gross margin (A6) gives the captured gross profit.

Step 3 — KPIs: Table IV summarises the comparison. Revenue and gross-profit figures are per vehicle per year; the fleet column scales the gross-profit delta linearly to a representative franchised group of 5,000 active customer vehicles.

TABLE IV
CURRENT-STATE / TARGET-STATE KPI COMPARISON. DELTA FIGURES
COMPUTED DIRECTLY ON UNROUNDED INTERMEDIATE VALUES; CURRENT
AND TARGET CELLS ROUNDED INDEPENDENTLY.

KPI	Current	Target	Δ
Repair capture rate	40 %	70 %	+30 pp
Wear-part revenue (/veh/yr)	€ 147	€ 258	+€ 110
Gross profit (/veh/yr)	€ 74	€ 129	+€ 55
Fleet GP, 5,000 veh (/yr)	€ 368k	€ 644k	+€ 276k
Customer touchpoints (/yr)	1.5	~2.5	+1.0
First-pass inspection rate	baseline	improved	qual.
Lifecycle retention [27]	~25 %	sustained	qual.

The proactive touchpoint count rises because each high-value forecast event (the wear parts in Table II excluding the wiper case, ≈ 1.0 events/yr) triggers an outreach contact that did not exist in the interval-based current state, roughly doubling the meaningful contact frequency from 1.5 to about 2.5 per year, assuming forecast contacts fall outside scheduled service visits; overlap would reduce the net gain modestly. Each additional contact is also an advisory and upselling opportunity; the parts targeted—brakes and tyres above all—are among the most common defect categories in the German mandatory vehicle safety inspection (HU) [29], so proactive replacement directly improves the customer’s first-pass inspection rate, a tangible and communicable benefit.

Sensitivity: The result depends most on the target capture rate. Varying c_{tgt} across a plausible band while holding V_{wear} and the 50 % margin fixed gives a per-vehicle gross-profit delta of:

$$c_{\text{tgt}} = 0.50 \rightarrow \Delta \text{ gross profit} \approx \text{€ } 18 \text{ /veh/yr,}$$

$$c_{\text{tgt}} = 0.60 \rightarrow \Delta \text{ gross profit} \approx \text{€ } 37 \text{ /veh/yr,}$$

$$c_{\text{tgt}} = 0.70 \rightarrow \Delta \text{ gross profit} \approx \text{€ } 55 \text{ /veh/yr,}$$

$$c_{\text{tgt}} = 0.80 \rightarrow \Delta \text{ gross profit} \approx \text{€ } 74 \text{ /veh/yr.}$$

Even at the conservative end the dealership captures additional gross profit on the order of tens of euros per vehicle per year, scaling to a six-figure annual sum across a mid-size fleet—before accounting for the compounding retention effect, namely that dealer-serviced customers are substantially more likely to purchase their next vehicle from the same dealership [31] (US data). The economic case is therefore robust to the assumptions rather than an artefact of any single optimistic value.

D. Limitations

The analysis is intentionally assumption-driven and carries corresponding limitations. First, the figures are illustrative: V_{wear} , the capture rates and the conversion factor are anchored to published market and accuracy data but are not fitted to an operational dataset, and no prototype was built; the contribution is a transparent model of the potential, not a measured outcome. Second, the PdM forecast accuracy (A9, 90 %) is reported for transparency but does not enter the capture-rate formula. Because a proactive service recommendation looks identical to the customer whether the underlying forecast is correct or slightly premature, both true-positive and false-positive predictions generate the same type of offer, accepted or declined at rate k . Accuracy will improve continuously as in-service data accumulates [9], but the model outcome does not depend on reaching any particular accuracy threshold. Third, and most consequential for practice, the target state presumes the dealership can actually obtain the vehicle’s telemetry. Under the EU Data Act (applicable from 12 September 2025) the user is the vehicle owner or driver, who may authorise a third party such as the dealership to receive in-vehicle data on fair, reasonable and non-discriminatory terms, while the OEM remains the data holder; where the data is personal, the GDPR additionally requires a lawful basis [32]. Access is thus contingent on customer consent, OEM provisioning in a usable format, and business-to-business (B2B) compensation — a non-technical adoption barrier that the model does not price in. Finally, two scope caveats apply: tyres, though included as a wear part, are often served by specialist channels and may leak independently of the dealership’s PdM capability; and the shift to electric vehicles will erode some ICE-specific wear-part volume (spark plugs, exhaust), shifting the addressable basket toward brakes, tyres and the high-voltage battery over the fleet’s lifetime.

V. DISCUSSION AND RECOMMENDATIONS

- Interpretation der Ergebnisse
- Implikationen für Forschung und Praxis

VI. CONCLUSION & FUTURE WORK

A. Conclusion

This paper examined the effect of digital-twin-based predictive maintenance on automotive aftersales, focusing on the German passenger-car market. A lifecycle comparison showed that the transition from interval-based servicing to a DT/PdM-enabled target state redirects wear-part work from the

independent aftermarket back to the franchised dealership by replacing reactive, fault-driven appointments with proactive, evidence-based service recommendations. A quantitative potential assessment, anchored to published German market data, translated this shift into concrete KPIs: under a conservative conversion assumption ($k = 0.50$), the repair capture rate rises from 40 % to 70 %, increasing gross profit by approximately €55 per vehicle per year and €276k annually across a 5000-vehicle fleet. Customer touchpoints grow from 1.5 to roughly 2.5 per year, each an additional advisory and upselling opportunity. The sensitivity analysis confirmed the economic case is robust across the full plausible parameter range. The dominant adoption barrier is non-technical: access to in-vehicle telemetry requires customer consent and OEM provisioning under the EU Data Act and GDPR, a hurdle any real-world deployment must resolve.

B. Future Work

Three directions stand out for future research. First, the assumption framework should be validated against an operational DT deployment, replacing modelled KPIs with measured outcomes and calibrating the conversion rate k under real customer-behaviour conditions. Second, the data-access challenge warrants dedicated study: viable consent architectures and OEM-dealer data-sharing agreements under the EU Data Act need to be developed and their effect on customer opt-in quantified. Third, the framework should be extended to battery electric vehicles, where the high-voltage battery's state of health becomes the dominant prognostic target and the addressable wear-part basket shifts away from ICE-specific components such as spark plugs.

ACKNOWLEDGMENT

The authors would like to thank Eduard Dobenko (Mercedes-Benz) for supervising this work.

REFERENCES

- [1] J. Lopes, A. Reis, L. Barbosa, and J. Barroso, "A vehicle digital twin in the context of a software ecosystem," in *2025 IEEE Vehicle Power and Propulsion Conference (VPPC)*, 2025, pp. 1–6.
- [2] R. Coppola and M. Morisio, "Connected car: Technologies, issues, future trends," *ACM Comput. Surv.*, vol. 49, no. 3, Oct. 2016. [Online]. Available: <https://doi.org/10.1145/2971482>
- [3] National Transport University, V. Biletskyi, O. Ivanushko, A. Loboda, O. Buhaichuk, National Transport University, R. Lovha, and "WebEye" (Ukraine), "Overview of information and communication technologies and vehicle monitoring systems and analysis of their possibilities for the formation and implementation of innovative technologies for maintenance and repair of vehicles," vol. 1, no. 273, pp. 8–16. [Online]. Available: http://journal.insat.org.ua/?page_id=4937&lang=en_GB
- [4] W. Alexander, J. Dempsey, S. Dayal, and J. Ark, "The secret life of factory service centers," p. 106.
- [5] G. Challagalla, R. Venkatesh, and A. K. Kohli, "Proactive postsales service: When and why does it pay off?" *Journal of Marketing*, vol. 73, no. 2, pp. 70–87, 2009. [Online]. Available: <https://doi.org/10.1509/jmkg.73.2.70>
- [6] O. Sabbagha, M. N. A. Rahman, W. R. Ismail, and W. M. H. W. Hussain, "Impact of quality management systems and after-sales key performance indicators on automotive industry: A literature review," *Procedia - Social and Behavioral Sciences*, vol. 224, pp. 68–75, 2016. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S1877042816304852>
- [7] M. Grieves, "Digital twin: Manufacturing excellence through virtual factory replication," *White Paper*, 2014.
- [8] C. Wang, I.-S. Fan, and S. King, "A review of digital twin for vehicle predictive maintenance system," in *Proceedings of the SAE AeroTech Conference*, 2023, sAE Technical Paper 2023-01-1024.
- [9] H. AbouGrad, V. Bathula, and F. Riaz, "Xai digital twin predictive maintenance in the automotive industry: Challenges and opportunities," in *Proceedings of the 2026 IEEE 5th International Maghreb Meeting of the Conference on Sciences and Techniques of Automatic Control and Computer Engineering (MI-STA)*, Sebha, Libya, 2026.
- [10] R. Wise and P. Baumgartner, "Go downstream: The new profit imperative in manufacturing," *Harvard business review*, vol. 28, pp. 89–96.
- [11] H. E. Hadraoui, N. Ouahabi, N. E. Bazi, O. Laayati, M. Zegrari, and A. Chebak, "Toward an intelligent diagnosis and prognostic health management system for autonomous electric vehicle powertrains: A novel distributed intelligent digital twin-based architecture," *IEEE Access*, vol. 12, pp. 110729–110761, 2024.
- [12] Karthick G, S. Devakatte, U. Kumar, D. Muthukumar, J. M. A. Sebastian, and J. P. Wagh, "Simulating the future: Ai-powered digital twins for intelligent emobility and software-defined vehicles," in *2025 IEEE International Transportation Electrification Conference (ITEC-India)*, 2025, pp. 1–4.
- [13] Vinay Kumar R, Pavan Kumar K. K., D. Ramasamy, M. M. Noor, W. S. W. Harun, L. Samylingam, N. Aslfattahi, M. S. Kiai, and C. K. Kok, "Digital twin framework for predictive maintenance of autonomous vehicle suspension system," in *2025 IEEE 4th International Conference on Smart Technologies for Power, Energy and Control (STPEC)*, 2025, pp. 1–6.
- [14] Y. Zhang, G. W. Gantt, M. J. Rychlinski, R. M. Edwards, J. J. Correia, and C. E. Wolf, "Connected vehicle diagnostics and prognostics, concept, and initial practice," *IEEE Transactions on Reliability*, vol. 58, no. 2, pp. 286–294, 2009.
- [15] S. Verstrepen, D. Deschoolmeester, and R. J. Berg, "Servitization in the automotive sector: creating value and competitive advantage through service after sales," in *Global Production Management*, K. Mertins, O. Krause, and B. Schallack, Eds. Boston, MA: Springer US, 1999, pp. 538–545. [Online]. Available: http://link.springer.com/10.1007/978-0-387-35569-6_66
- [16] M. W. M. Aljabar and S. Hasibuan, "Performance evaluation of after-sales service partners in the power tools industry," *Jurnal Sistem dan Manajemen Industri*, vol. 5, no. 2, pp. 105–114, 2021. [Online]. Available: <https://e-jurnal.lppmunsera.org/index.php/JSMI/article/view/3985>
- [17] A. G. González, "Service quality and repurchase behaviour in the spanish automotive after sales business," *Journal of Relationship Marketing*, vol. 14, no. 3, pp. 239–267, 2015. [Online]. Available: <http://www.tandfonline.com/doi/full/10.1080/15332667.2015.1069456>
- [18] J. R. Balinado, Y. T. Prasetyo, M. N. Young, S. F. Persada, B. A. Miraja, and A. A. N. Perwira Redi, "The effect of service quality on customer satisfaction in an automotive after-sales service," *Journal of Open Innovation: Technology, Market, and Complexity*, vol. 7, no. 2, p. 116, 2021. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S219985312200885X>
- [19] L. D. Peters, "IT enabled marketing: A framework for value creation in customer relationships," *Journal of Marketing Practice: Applied Marketing Science*, vol. 3, no. 4, pp. 213–229, 1997.
- [20] N. Saccani, L. Songini, and P. Gaiardelli, "The role and performance measurement of after-sales in the durable consumer goods industries: an empirical study," *International Journal of Productivity and Performance Management*, vol. 55, no. 3, pp. 259–283. [Online]. Available: <http://www.emerald.com/ijppm/article/55/3-4/259-283/165392>
- [21] R. Bundschuh and T. Dezvane, "How to make after-sales service pay off," *McKinsey Quarterly*, vol. 4, pp. 116–128.
- [22] A. Fraß, "Impact of after-sales performances of german automobile manufacturers in china in service satisfaction and loyalty," in *1st INTERNATIONAL CONFERENCE ON BUSINESS MANAGEMENT - "NEW CHALLENGES IN BUSINESS RESEARCH" - Conference proceedings*, Editorial Universitat Politècnica de València, Ed. Editorial Universitat Politècnica de València, pp. 200–206. [Online]. Available: <http://ocs.editorial.upv.es/index.php/ICBM/1ICBM/paper/view/1218>
- [23] M. Pophaley and R. K. Vyas, "Plant maintenance management practices in automobile industries: a retrospective and literature review," *Journal of Industrial Engineering and Management*, vol. 3, no. 3, pp. 512–541, 2010. [Online]. Available: <http://jiem.org/index.php/jiem/article/view/136>

- [24] V. González-Prida, C. P. Márquez, P. V. Gunckel, F. K. Rodríguez, and A. C. Márquez, "Digital transformation in aftersales and warranty management: A review of advanced technologies in I4.0," *Algorithms*, vol. 18, no. 4, p. 231, 2025. [Online]. Available: <https://www.mdpi.com/1999-4893/18/4/231>
- [25] D. F. Channon and T. Sammut-Bonnici, "Gap analysis," in *Wiley Encyclopedia of Management*, 1st ed., C. L. Cooper, Ed. Wiley, pp. 1–2. [Online]. Available: <https://onlinelibrary.wiley.com/doi/10.1002/9781118785317.weom120109>
- [26] Deutsche Automobil Treuhand GmbH, "Dat report 2025," DAT, Tech. Rep., Jan. 2025, analyse des Autojahres 2024. [Online]. Available: <https://www.dat.de/report/>
- [27] I. Maurer, N. Kramer, E. Scarinci, and M. Webster, "Optimizing dealer profitability with a service center tune-up," McKinsey & Company, Automotive & Assembly Insights, Nov. 2025, published Nov. 19, 2025, accessed June 3, 2026. [Online]. Available: <https://www.mckinsey.com/industries/automotive-and-assembly/our-insights/optimizing-dealer-profitability-with-a-service-center-tune-up>
- [28] Kraftfahrt-Bundesamt, "Verkehr in kilometern (vk) – inländer-fahrleistung 2024," https://www.kba.de/DE/Statistik/Kraftverkehr/VerkehrKilometer/vk_inlaenderfahrleistung/2024/2024_vk_kurzbericht.html, Jun. 2025, kurzbericht, abgerufen am 03.06.2026.
- [29] TÜV-Verband e.V., "TÜV-report 2026," TÜV-Verband e.V., Tech. Rep., Nov. 2025, basierend auf dem TÜV-Report 2026, abgerufen am 03.06.2026. [Online]. Available: <https://www.tuvsud.com/de-de/publikationen/tuev-report/die-haeufigsten-maengel>
- [30] T. Kroher, "Wartung, service, inspektion: Das müssen sie wissen," ADAC e.V., Jul. 2025. [Online]. Available: <https://www.adac.de/rund-ums-fahrzeug/reparatur-pflege-wartung/wartung-inspektion/wartungsintervalle/>
- [31] Cox Automotive, "2025 service industry study," Cox Automotive, Tech. Rep., Oct. 2025, survey of 1,974 U.S. vehicle owners conducted April–May 2025, accessed June 3, 2026. [Online]. Available: <https://www.coxautoinc.com/wp-content/uploads/2025/10/2025-Service-Industry-Study-FINAL.pdf>
- [32] European Parliament and Council of the European Union, "Regulation (eu) 2023/2854 on harmonised rules on fair access to and use of data (data act)," Official Journal of the European Union, Dec. 2023, oJ L, 22 December 2023, accessed June 3, 2026. [Online]. Available: <https://eur-lex.europa.eu/eli/reg/2023/2854/oj/eng>

VERWENDUNG GENERATIVER KI-SYSTEME

Bei der Erstellung der eingereichten Arbeit haben wir die nachfolgend aufgeführten auf künstlicher Intelligenz (KI) basierten Systeme benutzt:

- 1) Claude (Anthropic)
- 2)
- 3)

Wir erklären, dass wir

- uns aktiv über die Leistungsfähigkeit und Beschränkungen der oben genannten KI-Systeme informiert haben,
- die aus den oben angegebenen KI-Systemen direkt oder sinngemäß übernommenen Passagen gekennzeichnet haben,
- überprüft haben, dass die mithilfe der oben genannten KI-Systeme generierten und von uns übernommenen Inhalte faktisch richtig sind,
- uns bewusst sind, dass wir als Autorinnen bzw. Autoren dieser Arbeit die Verantwortung für die in ihr gemachten Angaben und Aussagen tragen.

Die oben genannten KI-Systeme haben wir wie im Folgenden dargestellt eingesetzt:

Tabelle I. Einsatz von KI-Systemen

Arbeitsschritt	Eingesetzte(s) KI-System(e)	Beschreibung der Verwendungsweise
Formatierung und Satz	Claude (Anthropic)	Umwandlung von Fließtext in $\text{\LaTeX}$ (Tabellen, Formeln, Literaturverzeichnis)
Verifikation der Berechnungen	Claude (Anthropic)	Prüfung der Zwischenrechnungen (V_{wear} , Capture-Raten) auf Konsistenz
Sprachliche Überarbeitung	Claude (Anthropic)	Korrektur von Grammatik und Stil einzelner Abschnitte